

Ekimetrics x Datacraft

Quand la baleine part dans les nuages

11/06/2021

datacraft^{*}
Ekimetrics.

Data science for business

PARIS | LONDON | NEW YORK | HONG KONG | DUBAI





Agenda.

- 00. Overview
- 01. Why Should I care ?
- 02. DE Overview
- 03. Let's go !
- 04. Training resources
- 04. Questions ?

What are we going to do ?

```
MyApp.py

import streamlit as st
import pandas as pd

st.write("""
# My first app
Hello *world!*
""")

df = pd.read_csv("my_data.csv")
st.line_chart(df)
```



Streamlit

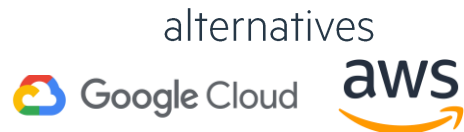
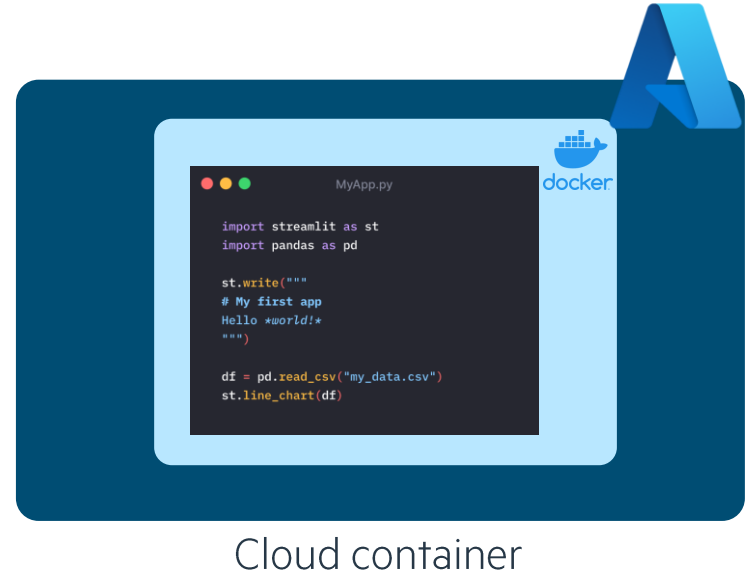
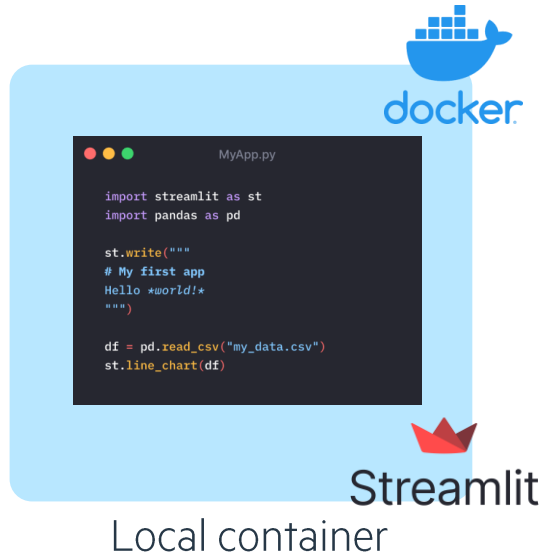


Local container

alternatives



What are we going to do ?



What do we need ?



Python



Github account



Dockerhub account



Vs Code



Azure Account

05. **Streamlit**

Streamlit : Data App made easy



<https://streamlit.io>

Created : October 2019

Python library (partly open-source)
Machine Learning & Data Science Apps

Code First & Low code

Keeping It Super Simple

```
MyApp.py

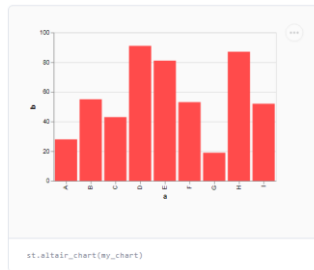
import streamlit as st
import pandas as pd

st.write("""
# My first app
Hello *world!*
""")

df = pd.read_csv("my_data.csv")
st.line_chart(df)
```

Interactive by design

Native sliders, box, calendar and more..



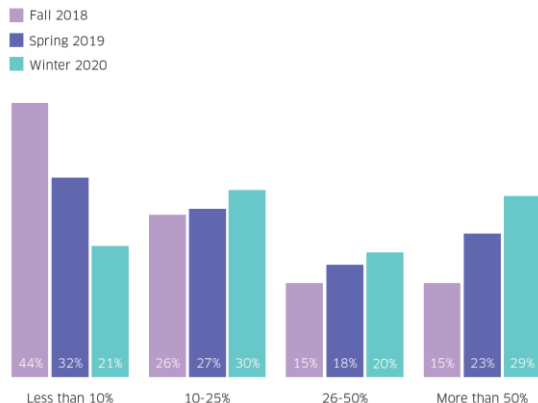
01.

Docker, why should I care ?

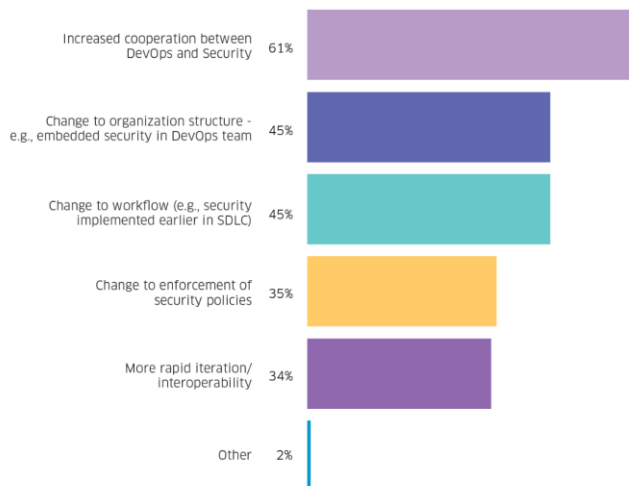
Everywhere

“By 2022, more than 75% of global organizations will be running containerized applications in production, up from less than 30% today.” – [Gartner, June 2020](#)

What percentage of your apps are containerized today?



How are containers changing how DevOps and Security work together?
(pick as many as apply)

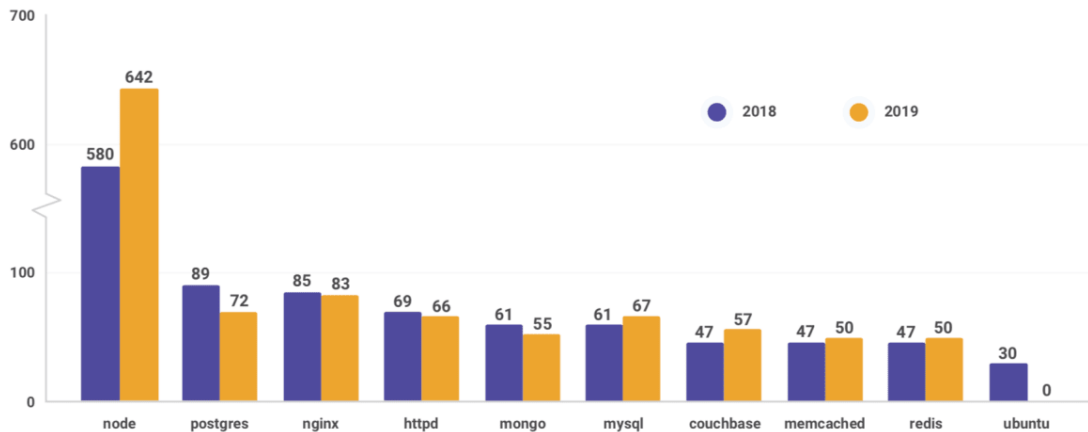


[source: stackrox](#)

What about security ?

Decrease of “native” vulnerabilities




Vulnerabilities in official container images



Rootless mode

introduced in Docker Engine v19.03 and graduated from experimental in Docker Engine v20.10. in Dec 2020.

Eki.Mission Example



Passed Security Audit

Why for Data Science ?

Reproducibility

Easily isolate and reproduce Data Science scenarios so that your colleagues obtain the same result on their end, whatever laptop they have.

Portability

No more friction when moving your PoC out of your laptop when more power is required. No need to recreate your env on the cloud, the container will handle it.

Scalability

With orchestration tools such as Swarm, Kubernetes, Kubeflow... you can scale your container to handle heavy loads.



The background of the slide is a soft-focus photograph of a desk. In the foreground, an open book with many pages is laid flat. Behind it, a copper-colored mug and a small potted plant with green leaves are visible, though they are out of focus.

01.

Some basic theory

Definitions



Dockerfile // The Recipe

Text file containing the set of instructions to build the docker image.

Docker build // The Cooking

Building an image from the Dockerfile.

Docker image // The Dough

A template file resulting of a docker build. If you run it, you create a container.

Docker run // The Baking

Run a command in a new container using a docker image.

Docker Container // The Cake

Set of isolated processes running on its own space on a shared kernel.

Docker client // You

Tool to interact with Docker.

Docker registry // Cake Shelf

Minimal collection of Docker layers

Docker hub // Cake Shop

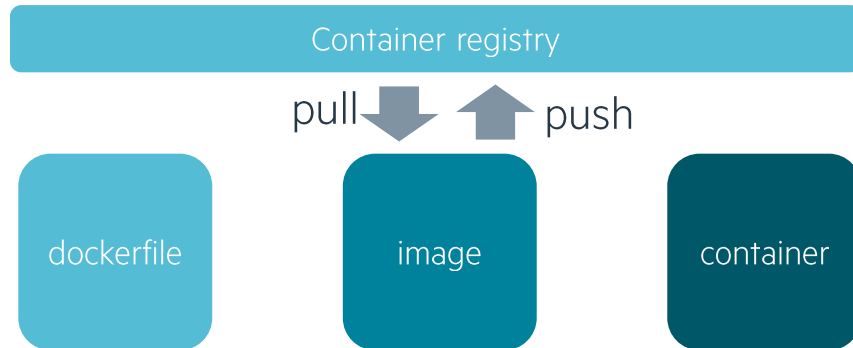
Centralized collection of Docker images, ready to be pulled. See it as GitHub for containers.

Docker pull // get a cake from the shelf

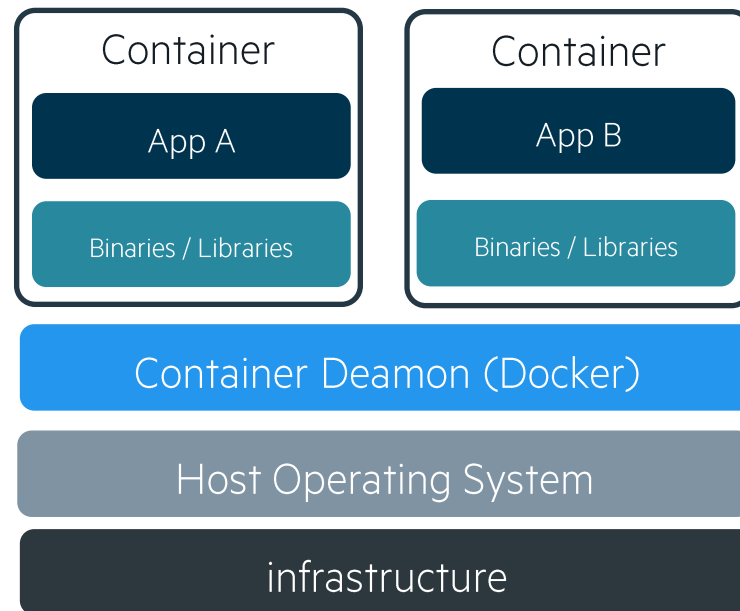
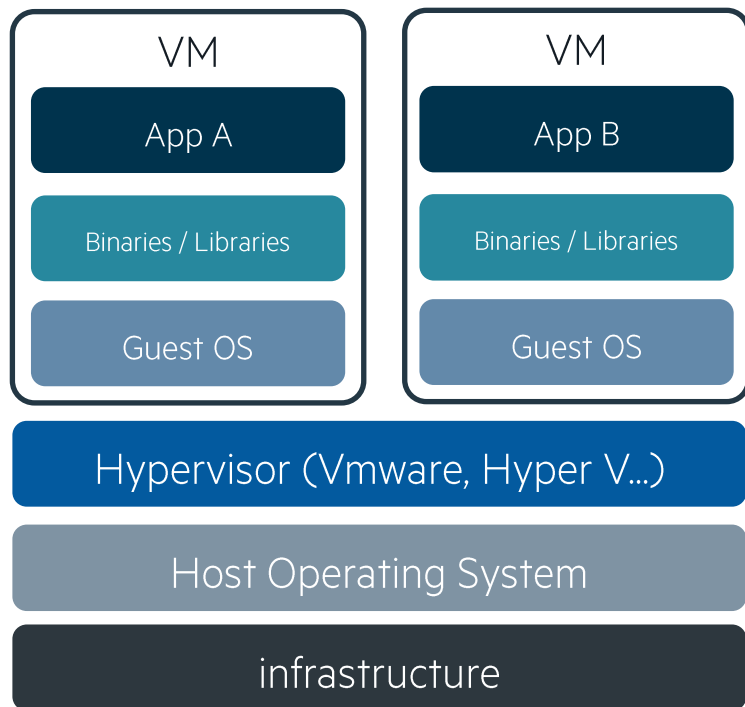
Like a git clone, get an image from a container registry

Docker push // put a cake on a shelf

Like a git push, send an image to a container registry



Container ≠ Virtual Machine



source: docker.com

Advantages of containers vs Virtual Machines

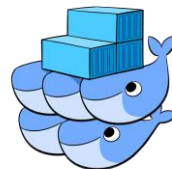


Portability

A container is like a self-sufficient package, with all its environment and dependencies shipped together as a single unit and can easily be moved around (on-premise or cloud)

Scalable

Containers can be “orchestrated” by tools such as Docker Swarm or Kubernetes that automate scaling, networking and deployment.



Maintanability

By separating every service in a container, you can easily update a container without impacting the other one

Simple example

```
docker run -dp 80:80 docker/getting-started
```

```
docker run -dp 80:80 hello-world
```

Common **Docker** commands

`docker image ls`

`docker image rm image_name:version`

`docker build -t image_name:version`

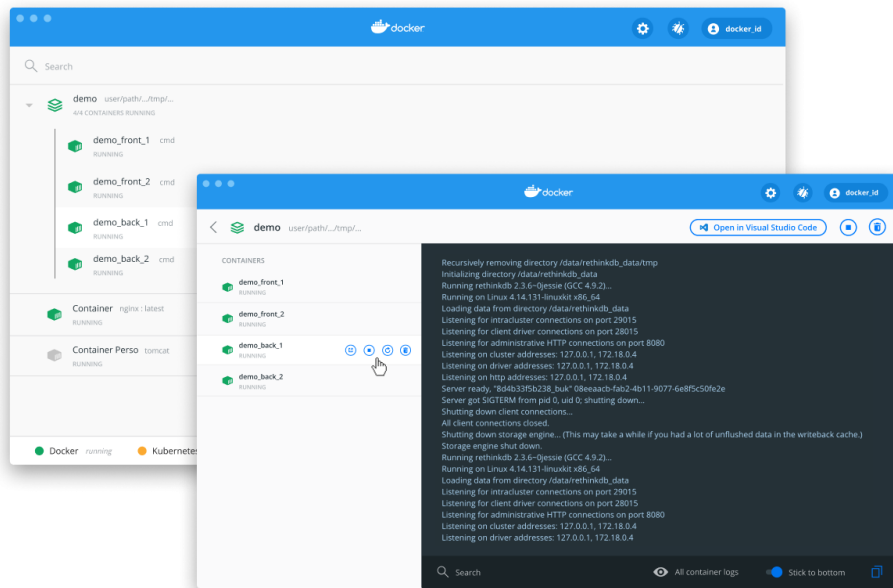
`docker container run --name container_name -p 80:80 image_name:version`

`docker pull image_name:version`

`docker tag image_name:version axelr/new_name:version`

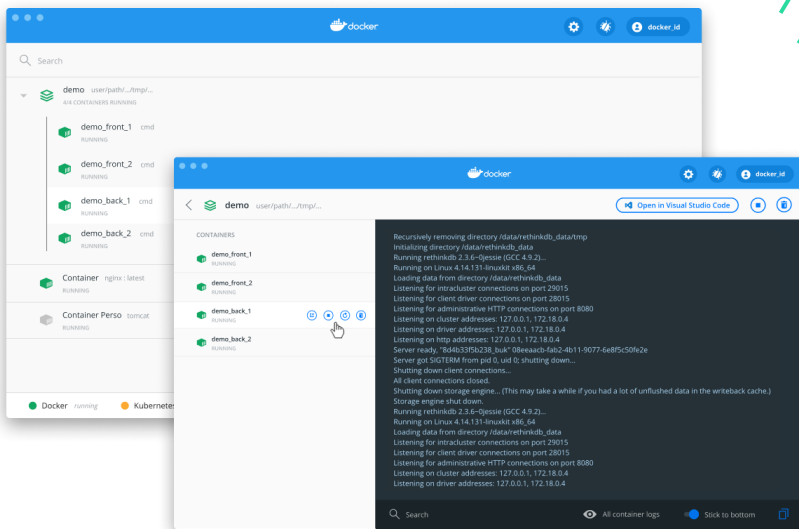
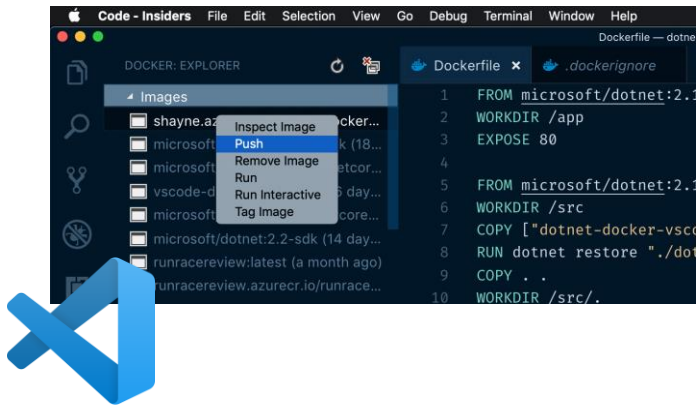
`docker push axelr/new_name:version`

Common Docker commands



Most commands are available with the Docker Desktop UI

Common Docker commands



Most commands are available within VS Code and/or the Docker Desktop UI

05.

Let's go



Getting started

- git clone [git@github.com:EatZeBaby/datacraft-workshop-docker.git](https://github.com/EatZeBaby/datacraft-workshop-docker.git)
- Start Docker Desktop
- code `./datacraft-workshop-docker`

Building our first Dockerfile

Choosing a base image is the first step of a Dockerfile = base linux distribution

```
FROM python:3.7
```

It's an image with python 3.7 preinstalled

Setting Working directory different from root

```
WORKDIR /app/
```

Installing the requirements for your app in a dedicated command (for caching purpose keep this method)

```
ADD ./requirements.txt /app/requirements.txt
```

```
RUN pip3 install -r requirements.txt
```

Copy everything necessary for the app to the image

```
ADD datacraft.py /app/datacraft.py
```

```
ADD css/ /app/css/
```

```
ADD utils/ /app/utils/
```

```
ADD models/ /app/models/
```

```
ADD images/ /app/images/
```

Run the app !

```
CMD ['streamlit', 'run', '/app/datacraft.py']
```

Build & Run your first container

- `docker build -t dockercraft .`
- `docker run -d -p 80:8501 dockercraft`
- Go to <http://localhost/>
- Update your code
- Refresh your web page
- What's wrong ? 😊

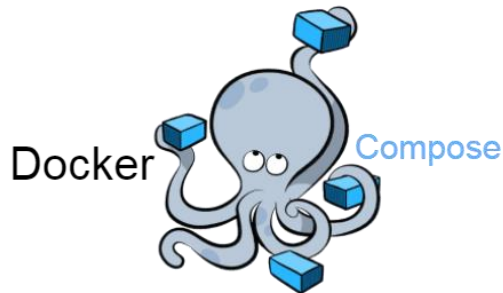
Docker Volumes

- The docker container is isolated from our source code.
- Docker volumes allows to create a link between local files and container files
- `docker volume create --name app_volum`
- `docker run -p 8050:8050`

Docker Compose

```
version: "3.8"
services:

  app:
    build:
      context: .
    container_name: "datacraft-docker"
    ports:
      - "8050:8050"
    volumes:
      - ./:/app
    env_file:
      - .env
```



- YAML file containing docker containers and multi container configuration.
- One off command « docker compose up » to start containers with many config options
 - port
 - env variables
 - Volumes
 -

Let's move onto the cloud

Push it to Docker Hub

- docker login
- docker tag dockercraft axreki/dockercraft:1
- docker push axreki/dockercraft:1

Azure set up

- Free Trial
- Create a resource group « datacraft-workshop-docker »
- [Quickstart center](#)
- [Create a container-based web app](#)
- Chose Free Tier for plan

Automating the deployment CI/CD

Automate build on Docker Hub = CI

git commit -am "commit message"

git tag -a v1.2 -m "version 1.2"

git push origin v1.2

Automated Builds

Autobuild triggers a new build with every **git push** to your source code repository. [Learn More.](#)

 [EatZeBaby/datacraft-workshop-docker](#) | Use Docker Hub's infrastructure | Autotests: Off

Docker Tag	Source	Latest Build Status	Autobuild	Build caching
version-{\1}	/^v([0-9.]+)\$/	SUCCESS	✓	✓

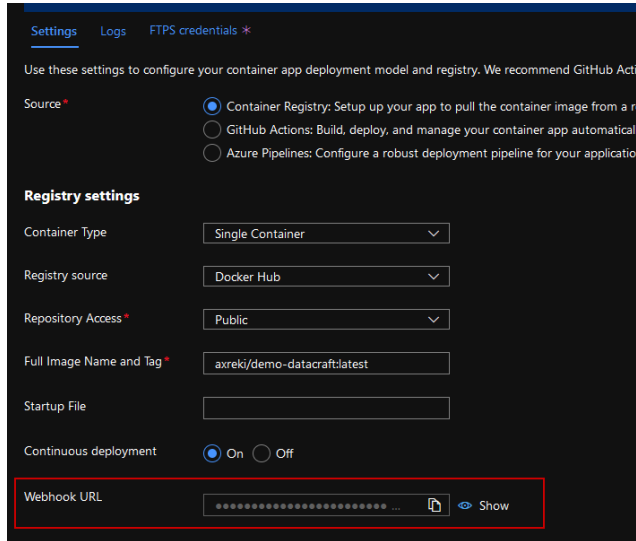
Automating the deployment CI/CD

Automate deployment on Azure = CD

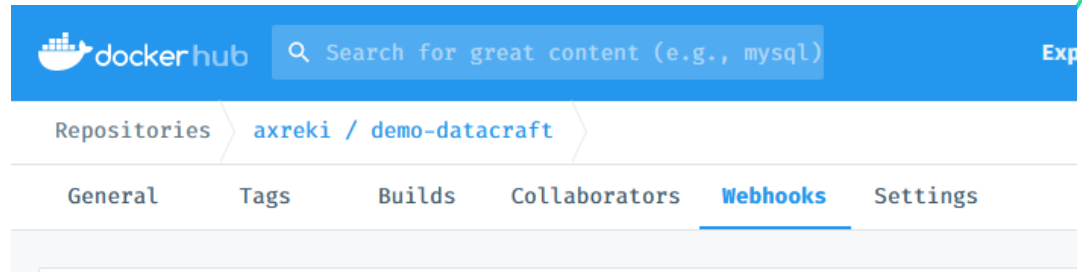
Get webhook url from Azure

Go to Azure > Web App > Container Settings > settings

Add a webhook to Docker Hub



The screenshot shows the 'Settings' tab in the Azure portal for a container app. Under 'Source', 'Container Registry' is selected. In the 'Registry settings' section, 'Container Type' is 'Single Container', 'Registry source' is 'Docker Hub', 'Repository Access' is 'Public', and 'Full Image Name and Tag' is 'axreki/demo-datacraft:latest'. 'Continuous deployment' is set to 'On'. The 'Webhook URL' field at the bottom is highlighted with a red box and contains a masked URL with a 'Show' button.



The screenshot shows the Docker Hub interface for the repository 'axreki / demo-datacraft'. The 'Webhooks' tab is selected in the top navigation bar. The page header includes the Docker Hub logo, a search bar, and the repository name. The 'Webhooks' tab is highlighted with a blue underline.

Automating the deployment CI/CD

Github Actions

Get webhook url from Azure

Go to Azure > Web App > Container Settings > settings

05. Optimization

Pro Tip: Reuse of top unmodified layer

Dockerfile

```
FROM python:3.7  
STEP 1.0  
STEP 2.0  
STEP 3.0  
STEP 4.0
```

First build

```
FROM python:3.7  
STEP 1.0  
STEP 2.0  
STEP 3.0  
STEP 4.0
```

Step 2 modified

```
FROM python:3.7  
STEP 1.0  
STEP 2.1  
STEP 3.0  
STEP 4.0
```

Step 4 modified

```
FROM python:3.7  
STEP 1.0  
STEP 2.0  
STEP 3.0  
STEP 4.0
```

build

reused cache

Pro Tip: Shared Layers

App A

```
FROM python:3.7
```

```
STEP 1.0
```

```
STEP 2.A
```

```
STEP 3.A
```

App B

```
FROM python:3.7
```

```
STEP 1.0
```

```
STEP 2.B
```

```
STEP 3.B
```

```
STEP 4.B
```

Reused cache

DÉCONSEILLÉ CAR 5 APPELS = 5 LAYERS

```
1  RUN apt-get update
2  RUN apt-get install -y apt-transport-https
3  RUN curl https://packages.microsoft.com/config/debian/9/prod.list > /etc/apt/sources.list.d/ms
4  RUN apt-get update
5  RUN apt-get install msodbcsql17 unixodbc-dev -y
```

1 SEUL LAYER

```
1  RUN apt-get update && \  
2      apt-get install -y apt-transport-https && \  
3      curl https://packages.microsoft.com/keys/microsoft.asc | apt-key add - && \  
4      curl https://packages.microsoft.com/config/debian/9/prod.list > /etc/apt/sources.list.d/ms  
5      apt-get update && \  
6      ACCEPT_EULA=Y apt-get install msodbcsql17 unixodbc-dev -y
```

EFFET BONUS : GARANTIE D'AVOIR UN APT-GET UPDATE TOUJOURS À JOUR

RÉDUCTION DOCKER SIZE

```
1 FROM python:3.7
2 WORKDIR /app
3 COPY Pipfile* ./
4 RUN pip install pipenv
5 RUN pipenv install --system --deploy
6 COPY src .
7 CMD ["python", "streamlit_index.py"]
```

1GO+

```
1 FROM python:3.7-slim
2
3 WORKDIR /app
4
5 # both files are explicitly required!
6 COPY Pipfile Pipfile.lock ./
7
8 RUN pip install pipenv && \
9     apt-get update && \
10     apt-get install -y --no-install-recommends gcc python3-dev libssl-dev && \
11     pipenv install --deploy --system && \
12     apt-get remove -y gcc python3-dev libssl-dev && \
13     apt-get autoremove -y && \
14     pip uninstall pipenv -y
15
16 COPY app ./
17
18 CMD ["python", "streamlit_index.py"]
```

~250mo

- Base image allégée (python-slim)
- unification des commandes RUN
- nettoyage des package uniquement nécessaire au build

Comparaison des tailles d'image python

```
axel.richier@EKI-PC00053 ~  
> docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
python	3.8	b4b9bf31ec03	23 hours ago	882MB
python	3.8.3-slim	9d84edf35a0a	6 months ago	165MB
python	3.8.3	7f5b6ccd03e9	6 months ago	934MB
python	3.8.3-alpine	8ecf5a48c789	6 months ago	78.9MB



A high-angle, slightly blurred photograph of a large number of empty black chairs arranged in rows in a hall. The chairs are modern, with a curved back and a simple metal frame. The floor is made of large, light-colored tiles. The lighting is soft, creating a calm and quiet atmosphere.

05.

Learn more

Learning resources

```
$> docker run -dp 80:80 docker/getting-started
```

[Play with Docker](#) [No install required]

[Containerized Python Development](#)

[3h YouTube Course](#)

Not Only Docker

[CoreOS rkt](#)

[Mesos Containerizer](#)

[LXC Linux Containers](#)

[OpenVZ](#)

[containerd](#)



Apache
MESOS



DANKE!
THANK YOU!
MERCİ!
GRAZIE!
GRACIAS!
DANK JE WEL!

.....

Questions ?

